

Sentiment Analysis: User Generated Comments on Streaming Platforms

Ganesh Khekare^{1*}, Shrey Gupta², Shreyansh Shailesh³, Goutam Majumder⁴, Parthiban Krishnamoorthy⁵

^{1,2,3,4,5}*School of Computer Science and Engineering, Vellore Institute of Technology, Vellore, India*

^{1*}*khekare.123@gmail.com*, ²*shrey.gupta2708@gmail.com*, ³*shreyansh.shailesh03@gmail.com*,
⁴*goutam.nita@gmail.com*, ⁵*k.parthiban@vit.ac.in*

Keywords: Natural Language Processing, Deep Learning, Transformer Models, Multimodal Sentiment Analysis, Text Classification

Abstract

Sentiment analysis, an important feature of natural language processing (NLP), focuses on rooting, relating, and grading the feelings, stations, opinions, or sentiments expressed in textbooks. As a decreasingly significant tool in the digital age, sentiment analysis has set up operations across different disciplines, including social media monitoring, political analysis, brand operation, client service, and request exploration. The capability to reuse and analyze large-scale textbook data has made sentiment analysis an inestimable fashion for gaining perception into public opinion, consumer behavior, and overall sentiment trends. This research provides an in-depth disquisition of sentiment analysis ways, pressing the colorful styles used to detect sentiment from textual data. The primary approaches bandied include traditional rule-grounded styles that work sentiment dictionaries, machine literacy ways like support vector machines and decision trees, and ultramodern deep literacy styles like Recurrent neural networks (RNN) and mills. These approaches, while important, face multitudinous challenges, like handling affront, nebulosity, contextual meaning, and the ever-growing complexity of mortal language. Also, the report covers the practical operations of sentiment analysis in real-world scripts. It examines how businesses use sentiment analysis for client feedback analysis, product review aggregation, and brand character operation. The report also addresses the limitations and future directions of sentiment analysis, including the need for more accurate models able to understand subtle nuances, multilingual sentiment discovery, and the integration of multimodal data sources (e.g., textbook, voice, and images) for further robust sentiment interpretation. Eventually, ethical considerations and impulses essential in sentiment analysis models are discussed, underlying the need for further transparent and fair algorithms.

1. Introduction

In the digital era, where a large amount of textual data is constantly being produced through social media, online reviews, blogs, forums, and news articles, the ability to automatically interpret and analyze the sentiment embedded in these texts has become an invaluable asset. Sentiment analysis or opinion mining is a process to examine the sentiment of the text, and the motive behind the sentiment analysis techniques is to judge the text based on emotion. The overall aim of the model is to understand the opinions of people around the world with the help of artificial intelligence like GPT and BERT. Sentiment Analysis has made huge progress, but the complexity of human languages has made it even more and more difficult for models to understand their emotions using AI tools like ChatGPT. It is easier to understand sarcasm and things that are not meant directly [1].

This report dives into the evolution of sentiment analysis over time. It describes the overall use of models like sentiment analysis in the daily world and how these types of advanced models help various industries and firms [2]. Sentiment Analysis, the task of determining the tone of the contextual datasets, has become a cornerstone in fields like marketing, customer service, and social media monitoring.

Thus, the problem is in developing a more robust and dependent model so that the potential of the model can be fully utilized, and it can handle diverse forms of expression and account for much multimodal information. The report delves into the evolution of sentiment models from scratch and takes note of all the deficiencies present in the current system models. It discusses the implications caused by the advancements of the models from textual data to voice recognition and facial recognition. These sentiment models have created an ease of doing judgment work, such as determining the context and mood of the data fed [3].

Despite being widely accepted as one of the key features in various industries, the challenge is to determine and analyze the mood and context of the data, whether it be positive, negative, or neutral, in more complex and big data sets in which the data available is vast [4]. Another significant challenge in these sentiment analysis models is the determination of textual data only. This limits the efficiency and potential of these models, limiting them to just textual datasets. To combine and form a model that recognizes all the aspects of the data, whether it be of any kind and nature. The uses of sentiment analysis have grown from time to time, and now is the time when these analysis models are used in almost every field whether it be marketing, finance, healthcare services, or political organizations making it a

very crucial model that needs to provide the user with actual sentiment records [5]. Thus, the problem is in creating a more robust and reliable model that could catch the sentiments of the data provided with good accuracy. A model that can analyze all types of data, whether it be textual data, auditory data, or facial recognition datasets.

2. Literature Review

Sentiment Analysis is a field of natural language processing (NLP) that focuses on extracting information from textual data, particularly from user-generated comments on platforms like YouTube, Twitch, etc. As sentiment models usually struggle with human text nature because of sarcasm, irony, etc., the introduction of models like BERT and GPT can help these models to cope with human text nature, such as slang [6]. The introduction of Deep Learning, Machine Learning, and Support Vector Machines add significantly to the potential of sentiment models and help them detect almost exact human nature and predict whether the context of the textual data is positive, negative, or neutral [7].

2.1 Early Approaches and Rule-Based Models

Early sentiment models were initially rule-based and were dependent on a predefined set of words and data. This type of approach is known as the lexicon-based approach. The most common form of sentiment analysis was to calculate a raw score, and based on the score, the data was judged as positive, negative, or neutral. For instance, [8] proposed a method based on the Mutual Information (PMI) score to identify the contexts and the mood of the data. His model worked well for short texts and words, but when it came to longer texts and more complicated texts, it was not considered effective.

2.2 Maintaining the Integrity of the Specifications

As sentiment models began to scale up, researchers started diving into machine learning algorithms to increase accuracy and adaptability. ML models automatically learned patterns, eliminating the need for lexicons. Pang and Lee [9] gave a review of some ML-based approaches, including techniques like Naive Bayes and Support Vector Machines. They explored the work of SVM for sentiment analysis and demonstrated that it is more efficient than the traditional methods used by researchers before. One of the major additions and progress in the ML domain was the addition of deep learning in the sentiment models, which helped in gaining much superior performance and even in more complex words and textual data [10]. This introduction gave better results than the traditional ML models, making them ideal for longer texts, blog posts, product reviews, and news articles.

2.3 Deep Learning and Transformer Models

In recent times, deep learning models, especially those on transformer architecture, have improved sentiment analysis tremendously. The transformer model introduced [11] marked a change from traditional models to model the connection between words in a sequence. This ability to

make use of sequences in parallel gave the chance to transformers to achieve state-of-the-art results across many NLP tasks. The release of models like BERT and GPT has increased the efficiency of sentiment models to much larger domains. BERT, which is a pre-trained transformer model on large-scale texts, fine-tuned it to some task-specific domains such as sentiment models. Following the success of BERT models like RoBERTa, they have started offering better performance in terms of the efficiency and accuracy of the sentiment models and datasets [12].

2.4 Multimodal Sentiment Analysis

While sentiment analysis models have been extensively studied and researched, new challenges and disparities are being discovered day by day, which is making analysis models work more difficult and complex. Now, when analyzing contextual data is fully under reach, researchers are focusing on visual and auditory features to enhance the analysis models. The proposed solution combines [13] all the textual data, visual data, and auditory data in one model. By integrating all these models into one integrated system, sentiment analysis models will be available to be used in more complex environments such as social media platforms and streaming platforms. Other works, such as that of [14], focused on integrating facial expression analysis into the multimodal system, which could already detect emotions in textual, auditorial, and visual data. This enhancement facilitated systems to detect human emotions using facial recognition of humans and helped in analyzing various video clips and findings.

2.5 Applications for Sentiment Analysis

Business and Marketing: Various companies use sentiment analysis to examine customer feedback and reviews, which indeed helps them to improve customer services.

Healthcare: Sentiment analysis is used in analyzing patients' health status using facial expressions and behavior. They also use sentiment analysis models to understand customer feedback and service requirements.

Political Analysis: Government and political organizations use sentiment analysis to track voters' sentiments and public opinion. They also track social movements and understand what the public demands using sentiment models. Sentimental models also help in making political decisions related to elections and public sentiments.

Crime Department: Organizations use sentiment analysis models in various crime department interrogations and hearings to determine whether the person being interrogated is providing the right information or not. The types of sentiment analysis models used in this kind of environment are facial recognition analysis and audio clip analysis.

2.6 Challenges in Sentiment Analysis

Algorithmic Bias: Biases in sentiment analysis models, such as gender, racial, and political biases, are a growing concern. Kiritchenko highlighted that certain models could reinforce stereotypes or provide skewed sentiment

classifications based on the demographic features of the authors of the text. There is ongoing research into developing fairness-aware sentiment models that address these biases [15].

Multilingual Sentiment Analysis: When we talk about sentiment analysis, we talk about a certain language, and the sentiment analysis model is based on a certain language, mostly English, which proves to be a challenge when the dataset provided is in a different language other than the language fed to the models. This can give wrong sentiment classifications and lead to flawed output [16].

Sarcasm and Irony: Sentiment analysis models often find it difficult to recognize instances of sarcasm and irony, which is when the actual meaning and the context of the data provided are different, or we can say, quite the opposite of what is fed in the system. This leads to wrong sentiment recognition, and thus, it is a challenge to be set right when it comes to providing the right sentiment classifications [17].

While significant progress has been made throughout the evolution of sentiment analysis from textual analysis to visual analysis, many areas remain unexplored and open for research. Many machine learning (ML) models, despite being very advanced and efficient, can show biases and lead to skewed data analysis when trained on a certain type of data. These biases can be related to gender, race, or political affiliation.

- **Multilingual and Cross-Cultural Sentiment Analysis:** Sentiment Analysis models are usually made to work under some language and cultural constraints, which is the reason why they fail when it comes to analyzing multilingual or cross-cultural datasets. While most of the sentiment analysis models have been made and trained in English, they face trouble and challenges when facing different languages having different syntax, idioms, and sentiment expressions. The lack of multilingual datasets and sentiment lexicons is the reason why this flaw still exists even after getting the limelight.

- **Bias in Sentiment Analysis Systems:** The most pressing concern in research is the existence of algorithm bias. Sentiment analysis models, particularly those based on machine research, have focused on batch-processing sentiment analysis. There is still a lack of efficient models capable of processing large-scale, real-time data promptly. Optimizing sentiment analysis algorithms for real-time applications as customer service chatbots, live sentiment tracking, or social media monitoring remains an open research area that has the potential to drive the next wave of innovation in sentiment analysis.

- **Sentiment Shifting and Temporal Analysis:** Temporal Analysis faces problems in a dynamic environment where the sentiment tends to change with time. In dynamic settings such as social media platforms or public sentiments, the sentiment changes with ongoing events, which results in flawed sentiment analysis. Another case where sentiment analysis fails to deliver is in a long-term context, for example, the sentiments before and after elections, product launches, or global crises.

- **Real-Time Sentiment Analysis:** Finally, the ability to perform real-time sentiment analysis on larger volumes of constantly updating data presents a significant technical challenge. While most of the research has focused on batch-processing sentiment analysis, there is still a lack of efficient models capable of processing large scale real-time data on time.

- **Domain Specific Sentiment Analysis:** While most sentiment analysis models perform well in general purpose tasks, they fail to deliver in domain specific tasks where the stress is on a specific task and context. It generalizes the data provided and cannot focus on the specific domain under which it has to work. Domain Adaptation Techniques are needed to fulfil the needs of these domain specific sentiment models so that they can provide accurate and valuable output to the user.

- **Contextual and Subtle Sentiment Decisions:** A major form of flaw we can say is the detection of textual data containing hidden meaning like data containing sarcasm and irony. Sentiment analysis models are unable to detect this kind of context without the help of AI. Existing models often misunderstand the context of the data. Things like sarcasm and irony used as the actual meaning of the data are the opposite of what is provided in the datasets.

3. Methodology

The sentiment analysis system development process involves several stages, starting with data collection. Text data was gathered from Kaggle in the form of a labeled dataset, which consisted of 10,000 labeled reviews, each classified as positive or negative. After the data collection phase, the data goes through the pre-processing phase. In this phase, there are various steps through which the data is passed. In the first step, all the data is converted into lowercase to standardize the data, ensuring the removal of case-sensitive deviations. The next step is noise removal, which means removing data that is not used when classifying sentiment. This includes special characters, stop words, URLs, and punctuation marks. The third step is tokenization, which means breaking the text down into smaller units. The tools used for this were NLTK and spaCy. It generates a vocabulary of terms, with each unique term mapped to a specific column. The next step is stemming. This is used to standardize the generated vocabulary. Additionally, the text data is vectorized by incorporating the help of tools like Word2Vec and TF-IDF.

Once the preprocessing steps are complete, we start the training and testing phase, in which the data is then split into training and testing in a ratio of 4 to 1. That means 80% of the data is used for training the models, while 20% is used for testing and evaluation. Testing makes sure the model is effective under various circumstances. To further enhance the model's performance, MinMaxScaler was incorporated to normalize the term frequencies. Normalization makes sure that all features are consistently scaled in a range of 0 to 1. This step is crucial as it ensures that a particular raw feature with a large value doesn't skew the model's performance. It ensures that all features are considered

equal. In the end, the model's performance was evaluated using metrics like accuracy, precision, recall, and F1-score. section here.

4. Results and Discussion

The three Naive Bayes models, GaussianNB, MultinomialNB, and BernoulliNB, were tested on the test dataset, and their performance was evaluated using accuracy, precision, recall, and F1-score metrics. The accuracy scores for each model were as follows: GaussianNB scored 77.45%, MultinomialNB scored 82.85%, and BernoulliNB scored 82.50%. The same is shown in Figure 1.




Figure 1 Accuracy Scores

The performance differences between these models highlight the importance of choosing the appropriate classifier for specific types of data. MultinomialNB outperformed the other two models, achieving the highest accuracy of 82.85%. This is because Multinomial Naive Bayes is predominantly employed for text classification tasks, where the features are represented as discrete term frequencies. It effectively models the frequency of occurrence of terms in text data, which is a natural fit for sentiment analysis tasks. GaussianNB, on the other hand, underperformed with an accuracy of only 77.45%. This can be attributed to the fact that Gaussian Naive Bayes assumes that features (such as word frequencies) follow a continuous distribution. However, text data is inherently discrete, as it consists of counts or occurrences of words, making the Gaussian distribution assumption less compatible with this type of data. As a result, GaussianNB did not perform as well in this context.

BernoulliNB, which assumes binary features (presence or absence of words), achieved an accuracy of 82.50%, slightly lower than MultinomialNB. While it performed comparably, it is more suited to datasets where features are binary, such as spam detection or when modeling the presence or absence of specific terms in a document. The model was not as effective when it came to longer reviews, where the frequency of terms becomes important when it comes to classifying sentiment, which is the reason for it being less effective as compared to MultinomialNB. The model was effective when it came to classifying short text.

Comparative Analysis: MultinomialNB achieved higher results in both metrics, precision, and recall, making it the most suited for tasks that require a balanced approach when classifying positive and negative reviews. Hence, it was chosen for deployment, as it regularly got the better of other models in the above-mentioned metrics. The deployed system was tested using real-time input, and it classified the reviews as intended. For example, "superb movie" was classified as positive, and "bad film" was classified as negative, as shown in Figure 2 and Figure 3.




Figure 2 Positive Classification




Figure 3 Negative Classification

The capability to deploy such a system with real-time input is important for practical applications like analyzing product reviews, customer feedback, social media posts, and more. Businesses can use sentiment analysis to tailor their approach and strategies to enhance feasibility and customer satisfaction. However, the system has a few limitations as well. These limitations include handling sarcasm, irony, and extremely long reviews. Another limitation is the use of multiple languages in the same sentence. The reviewer might write most of the sentence in English but might write some part of it in Hinglish, which means writing Hindi words using the English alphabet that sounds similar in pronunciation. These challenges can be overcome by future advancements, such as using advanced models like deep learning models or using additional features like context-aware embeddings.

5. Conclusion

The research displays the productivity of various Naive Bayes classifiers. MultinomialNB was especially effective for text-based classification. Achieving 83.2% accuracy, the model handled discrete term frequencies well, making it a preferred choice for structured textual data. Compared to other Naive Bayes classifiers (GaussianNB and BernoulliNB), MultinomialNB was better in metrics like accuracy, precision, and recall, proving its efficacy for this task. Despite the higher metrics, there were a few limitations. The inability to handle nuances like irony, sarcasm, and mixed emotions was one of the most obvious ones. These variations are usually difficult to identify using traditional algorithms. For example, a review like "I liked the movie...not" might be wrongly classified as positive because of the word "loved", even though to a human reader, it is negative. This shows that better models, such as deep learning models, need to be developed to overcome these limitations. Limited vocabulary was another limitation of this study. The model was trained by employing a vocabulary of 1000 words, which was done to reduce the dimensionality of the data. However, it can lead to the model missing out on various important words, particularly in longer data points. Certain words that can be domain specific might also not be part of the vocabulary used. This can weaken the model's intended usage and lead to degraded performance when dealing with more complex data, as it has limited data to train itself. Another limitation is the fact that although Naive Bayes is productive for most of the text classification tasks, it relies on the assumption that the various features are independent of each other. This fact may not hold in more sophisticated language data. This can further degrade the model's performance when dealing with highly interdependent data.

In conclusion, the model is effective in classifying reviews as positive or negative, with MultinomialNB proving to be the most effective model, scoring high in various metrics (precision, accuracy, recall). This study holds up the potential of Naive Bayes as a baseline model for sentiment analysis. However, it is limited when it comes to handling linguistic nuances like sarcasm, irony, negation, and mixed emotions. It is also limited due to the limited vocabulary size. Therefore, with further advancements, sentiment analysis can have large-scale applications in sectors like business, finance, marketing, healthcare, and more.

6. Future Scope

Methodologies: Research other techniques used in sentiment analysis, employing machine learning, deep learning, and more. Also exploring recent growth in transformer models such as BERT, GPT, and their ability to identify context.

Applications: Application of sentiment analysis in sectors like business, politics, finance, and healthcare. Other sectors like social media, marketing, and brand management can also employ sentiment analysis.

Challenges and Limitations: Improvements can be made to overcome challenges like detecting sarcasm and irony.

Other challenges, like ambiguous expressions and sentences that use words from multiple languages, also need to be overcome as the field is explored more with time.

Future Directions: Potential progress in sentiment analysis can be made by integrating multimodal data, which includes text, audio, and images. This can improve algorithmic efficiency and can be used to develop better systems that can interpret complex human emotions.

Challenges and Limitations: Improvements can be made to overcome challenges like detecting sarcasm and irony. Other challenges, like ambiguous expressions and sentences that use words from multiple languages, also need to be overcome as the field is explored more with time

7. References

- [1] Pang, B., & Lee, L. (2008). Opinion mining and sentiment analysis. *Foundations and Trends in Information Retrieval*, 2(1-2), 1-135.
- [2] Liu, B. (2012). Sentiment analysis and opinion mining. *Synthesis Lectures on Human Language Technologies*, 5(1), 1-167.
- [3] VaderSentiment (2014). VADER: A Parsimonious Rule-based Model for Sentiment Analysis of Social Media Text. *Proceedings of the 8th International Conference on Weblogs and Social Media*.
- [4] G. Khekare, C. Masudi, Y. K. Chukka and D. P. Koyyada, "Text Normalization and Summarization Using Advanced Natural Language Processing," *2024 International Conference on Integrated Circuits and Communication Systems (ICICACS)*, Raichur, India, 2024, pp. 1-6, doi: 10.1109/ICICACS60521.2024.10498983.
- [5] Cambria, E., & White, B. (2014). Jumping NLP curves: A review of natural language processing research. *IEEE Computational Intelligence Magazine*, 9(2), 48-57.
- [6] Khekare, Ganesh, Pushpneel Verma, and Seema Raut. "The Smart Accident Predictor System using Internet of Things." In *Cloud IoT*, pp. 163-175. Chapman and Hall/CRC, 2022.
- [7] Socher, R., Perelygin, A., Wu, J., Chuang, J., Manning, C. D., Ng, A. Y., & Potts, C. (2013). Recursive deep models for semantic compositionality over a sentiment treebank. *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing (EMNLP 2013)*, 1631-1642.
- [8] Bollen, J., Mao, H., & Zeng, X. (2011). Twitter mood predicts the stock market. *Journal of Computational Science*, 2(1), 1-8.
- [9] Khekare, G., Bramhane, L.K., Dhule, C., Agrawal, R., Turukmane, A.V. (2022). Testing and Analysis of Predictive Capabilities of Machine Learning Algorithms. In: Houssein, E.H., Abd Elaziz, M., Oliva,

- D., Abualigah, L. (eds) Integrating Meta-Heuristics and Machine Learning for Real-World Optimization Problems. Studies in Computational Intelligence, vol 1038. Springer, Cham. https://doi.org/10.1007/978-3-030-99079-4_16
- [10] Zhang, L., & Liu, B. (2018). Deep learning for sentiment analysis: A survey. Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery, 8(4), e1262.
- [11] Zhang, Y., & Zhu, J. (2019). Aspect-based sentiment analysis: A survey. IEEE Transactions on Affective Computing, 10(3), 338-348.
- [12] Poria, S., Hazarika, D., Majumder, N., & Gelbukh, A. (2017). A deeper look into sarcasm detection in text: A survey. ACM Computing Surveys (CSUR), 50(5), 1-22.
- [13] Joulin, A., Grave, E., Mikolov, T., & Bach, F. (2017). Bag of Tricks for Efficient Text Classification. arXiv:1607.01759.
- [14] Khekare, G., Turukmane, A.V., Dhule, C., Sharma, P., Kumar Bramhane, L. (2022). Experimental Performance Analysis of Machine Learning Algorithms. In: Qian, Z., Jabbar, M., Li, X. (eds) Proceeding of 2021 International Conference on Wireless Communications, Networking and Applications. WCNA 2021. Lecture Notes in Electrical Engineering. Springer, Singapore. https://doi.org/10.1007/978-981-19-2456-9_104
- [15] S. Zhao *et al.*, "SDRS: Sentiment-Aware Disentangled Representation Shifting for Multimodal Sentiment Analysis," in *IEEE Transactions on Affective Computing*, doi: 10.1109/TAFFC.2025.3539225.
- [16] G. Khekare, G. Majumder, N. Shelke, G. Sakarkar and S. Buchade, "A Deep Dive into Existing Lip Reading Technologies," *2024 International Conference on Artificial Intelligence and Quantum Computation-Based Sensor Application (ICAIQSA)*, Nagpur, India, 2024, pp. 1-6, doi: 10.1109/ICAIQSA64000.2024.10882250.
- [17] A. Jannani, S. Bouhsissin, N. Sael and F. Benabbou, "Topic Modeling and Sentiment Analysis of Arabic News Headlines for a Societal Well-Being Scoring and Monitoring System: Moroccan Use Case," in *IEEE Access*, doi: 10.1109/ACCESS.2025.3538888.